

Computer Networks

EDA387/DIT663

Fault-tolerant Algorithms for Computer Networks

Super-stabilization (Ch. 7)

Chapter 7: roadmap

7.1 Superstabilization

7.2 Self-Stabilizing Fault-Containing Algorithms

Dynamic System & Self Stabilization

Dynamic System

Algorithms for *dynamic systems* are designed to cope with failures of processors with no global re-initialization.

Such algorithms consider only global states reachable from a predefined initial state under a restrictive sequence of failures and attempt to cope with such failures with as few adjustments as possible.

Self Stabilization

Self-stabilizing algorithms are designed to guarantee a particular behavior finally.

Traditionally, changes in the communications graph were ignored.



Superstabilizing algorithms combine the benefits of both self-stabilizing and dynamic algorithms

Definitions

A Superstabilizing Algorithm:

- Must be self-stabilizing
- Must preserve a “*passage predicate*”
- Should exhibit fast convergence rate

Passage Predicate - Defined with respect to a class of topology changes (A topology change falsifies legitimacy and therefore the passage predicate must be weaker than legitimacy but strong enough to be useful).

Passage Predicate - Example

In a token ring:

Passage Predicate	Legitimate State
<u>At most</u> one token exists in the system. (e.g. the existence of 2 tokens isn't legal)	<u>Exactly</u> one token exists in the system.

A processor crash can lose the token but still not falsify the passage predicate

Evaluation of a Super-Stabilizing Algorithm

a. Time complexity

The maximal number of rounds that have passed from a legitimate state through a single topology change and ends in a legitimate state

b. Adjustment measure

The maximal number of processors that must change their local state upon a topology change, in order to achieve legitimacy

Motivation for Super-Stabilization

A self-stabilizing algorithm that does not ignore the occurrence of topology changes (“events”) will be initialized in a predefined way and react better to dynamic changes during execution

Question:

Is it possible, for the algorithm that detects a fault, when it occurs, to maintain a “*nearly legitimate*” state during convergence?

Motivation for Super-Stabilization



While transient faults are rare (but harmful), a dynamic change in the topology may be frequent.

Thus, a super-stabilizing algorithm has a lower worst-case time measure for reaching a legitimate state again, once a topology change occurs.

In the following slides we present a self-stabilizing and a super-stabilizing algorithm for the graph coloring task.

Graph Coloring

- a. The coloring task is to assign a color value to each processor, such that no two neighboring processors are assigned the same color.
- b. Minimization of the colors number is not required. The algorithm uses $\Delta+1$ colors, where Δ is an upper bound on a processor's number of neighbors.
- c. For example:



Graph Coloring - A Self-Stabilizing Algorithm

01 Do forever

02 $GColors := \emptyset$

03 For $m:=1$ to δ do

04 $lr_m := \text{read}(r_m)$

05 If $ID(m) > i$ then

06 $GColors := GColors \cup \{lr_m.\text{color}\}$

07 od

08 If $\text{color}_i \in GColors$ then

09 $\text{color}_i := \text{choose}(\setminus GColors)$

10 Write $r_i.\text{color} := \text{color}_i$


11 od

Colors of P_i 's

neighbors

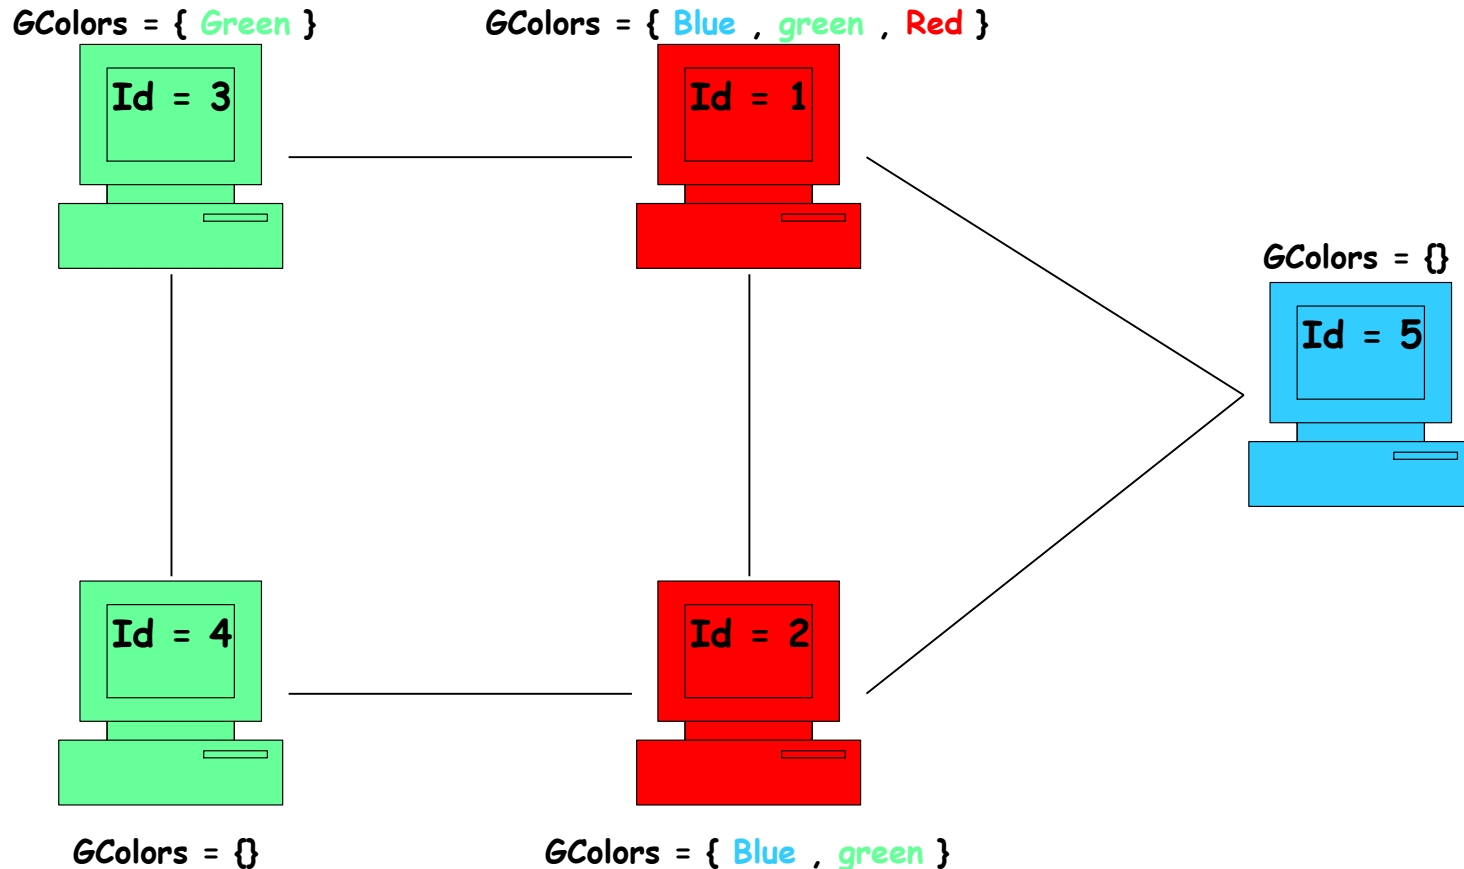
If P_i has the color of one of its neighbors with a higher ID, it chooses another color and writes it.

Graph Coloring - Self-Stabilizing Algorithm - Simulation



Phase I


Graph Coloring - Self-Stabilizing Algorithm - Simulation





Phase II

Graph Coloring - Self-Stabilizing Algorithm - Simulation


$GColors = \{ \text{Green} \}$



$GColors = \{ \text{Blue}, \text{green}, \text{Red} \}$




$GColors = \{ \}$



$GColors = \{ \text{Blue}, \text{green} \}$

$GColors = \{ \}$



Stabilized

Phase III

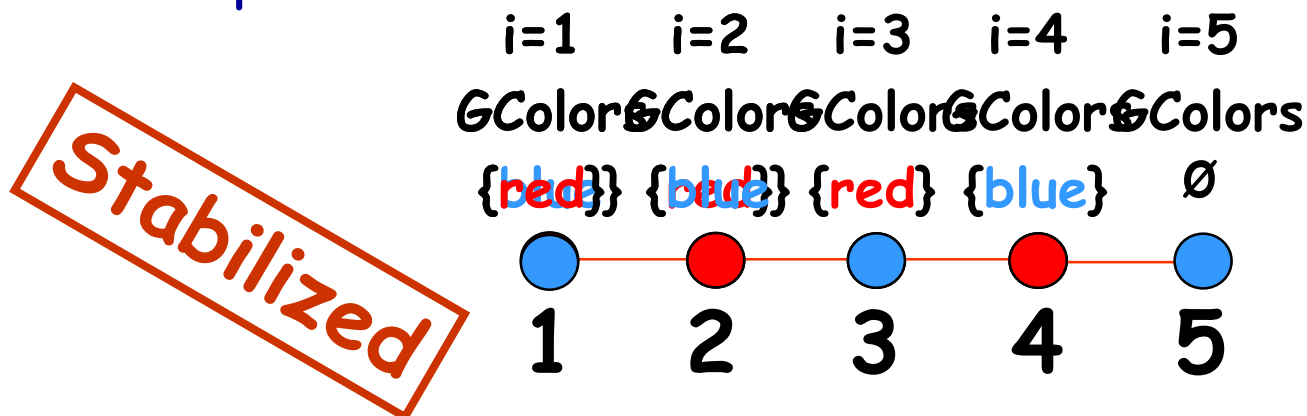
Graph Coloring - Self-Stabilizing Algorithm (continued)

What happens when a change in the topology occurs ?

If a new neighbor is added, it is possible that two processors have the same color.

It is possible that during convergence every processor will change its color.

Example:



But in what cost ?

Graph Coloring - Super-Stabilizing Motivation

- a. Every processor changed its color but only one processor really needed to.
- b. If we could identify the topology change we could maintain the changes in its environment.
- c. We'll add some elements to the algorithm:
 - a. AColor - A variable that collects all of the processor neighbors' colors.
 - b. Interrupt section - Identify the problematic area.
 - c. \perp - A symbol to flag a non-existing color.

Graph Coloring - A Super-Stabilizing Algorithm

```
01  Do forever
02      AColors :=  $\emptyset$ 
03      GColors :=  $\emptyset$ 
04      For m:=1 to  $\delta$  do
05          lrm:=read(rm)
06          AColors:=AColors U {lrm.color}
07          If ID(m)>i then GColors := GColors U {lrm.color}
08      od
09      If colori =  $\perp$  or colori  $\in$  GColors then
10          colori:=choose( $\setminus \setminus$  AColors)
11      Write ri.color := color
12  od
13  Interrupt section
14      If recoverij and j > i then
15          Colori :=  $\perp$ 
16          Write ri.color :=  $\perp$ 
```

All of P_i neighbors' colors

Activated after a topology change to identify the critical processor

recover_{ij} is the interrupt which P_i gets upon a change in the communication between P_i and P_j

Graph Coloring - Super-Stabilizing Algorithm - Example

Note that the new algorithm stabilizes faster than the previous one (in some cases).

Let us consider the previous example, this time using the super-stabilizing algorithm:

$$\text{Color}_4 = \perp$$


$$r_4.\text{color} = \perp$$

$$i=4$$

$$G\text{Colors} = \{\text{blue}\}$$

$$A\text{Colors} = \{\text{blue}, \text{red}\}$$

Stabilized



In $O(1)$.

Graph Coloring - Super-Stabilizing Proof

Lemma 1: This algorithm is self-stabilizing.

Proof by induction:

a. After the first iteration:

a. The value \perp doesn't exist in the system.

b. P_n has a fixed value.


b. Assume that P_k has a fixed value $\forall i < k \leq n$.

If P_i has a P_k neighbor then P_i does not change to P_k 's color, but chooses a different color.

Due to the assumptions we get that P_i 's color becomes fixed for $1 \leq i \leq n$, so the system stabilizes.

Graph Coloring - Super-Stabilizing

Passage Predicate - The color of a neighboring processor is always different in every execution that starts in a safe configuration, in which only a single topology change occurs before the next safe configuration is reached



Graph Coloring - Super-Stabilizing

Super-stabilizing Time - Number of cycles required to reach a safe configuration following a topology change.

Super-stabilizing vs. Self-Stabilizing


$O(1)$

$O(n)$

Graph Coloring - Super-Stabilizing

Adjustment Measure - The number of processors that changes color upon a topology change.

The super-stabilizing algorithm changes one processor color, the one which had the single topology change



Reading Review

- Shlomi Dolev, Ted Herman: *Superstabilizing Protocols for Dynamic Distributed Systems*. Chicago J. Theor. Comput. Sci. 1997 (1997)

<http://cjtc.cs.uchicago.edu/articles/1997/4/cj97-04.pdf>